

Chapter 3

Implementing Neural Networks

Content of this Chapter

1. Neural Network Libraries
2. Introduction to PyTorch
 1. From Numpy to PyTorch
 2. Implementing a Neural Network
3. Vectorisation of Neural Networks


3.1 Neural Network Libraries

- What libraries exist?
- What are their features?
- Why is GPU support so important?
- What do we use here?

Deep Learning Package Zoo

- PyTorch
- Tensorflow
- Keras
- JAX
- ...

PYTORCH



Deep Learning Package Design Choices

- Model specification?
- Computational graph?
- High-level programming language?

Deep Learning Package Design Choices

Model specification

```
name: "convolution"
input: "data"
input_dim: 1
input_dim: 1
input_dim: 100
input_dim: 100
layer {
  name: "conv"
  type: "Convolution"
  bottom: "data"
  top: "conv"
  convolution_param {
    num_output: 3
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "gaussian"
      std: 0.01
    }
    bias_filler {
      type: "constant"
      value: 0
    }
  }
}
```

Configuration file

e.g.

- Caffe
- DistBelief
- CNTK

Programmatic generation

e.g.

- PyTorch
- Theano
- Tensorflow

```
import torch.nn as nn

class NewsgroupsModel(nn.Module):
    """Simple Feedforward Neural Network for 20
    Newsgroups"""
    def __init__(self, input_size=300):
        super().__init__()

        self.input_size = input_size
        self.hidden_1_size = 2048
        self.hidden_2_size = 256
        self.num_classes = 20

        self.fc1 = nn.Linear(self.input_size,
                              self.hidden_1_size)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(self.hidden_1_size,
                              self.hidden_2_size)
        self.relu2 = nn.ReLU()
        self.fc3 = nn.Linear(self.hidden_2_size,
                              self.num_classes)

    def forward(self, x):
        a = self.relu1(self.fc1(x))
        b = self.relu2(self.fc2(a))
        c = self.fc3(b) # => logits

        return c
```

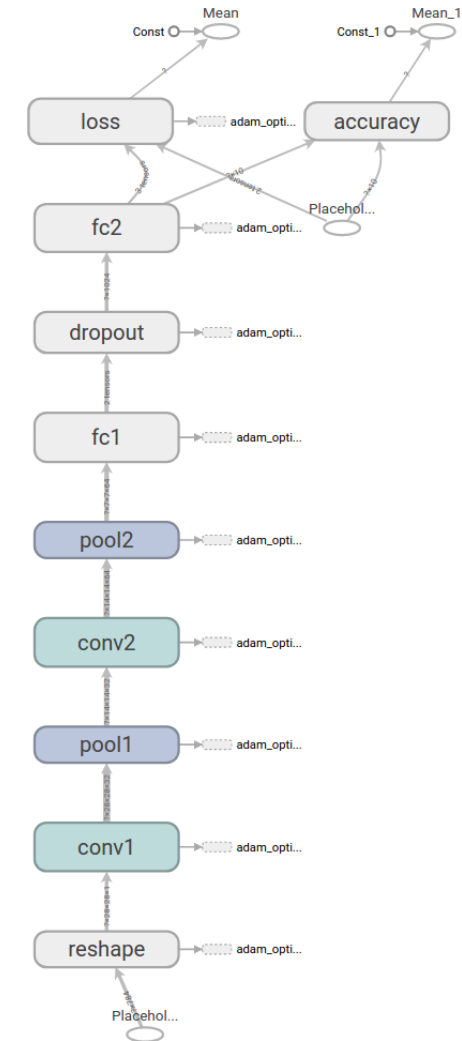
Deep Learning Package Design Choices

Computational graph

computational graph of a convolutional neural network:

- contains the program's operations and variables as a directed acyclic graph
- looks a lot like the circuit diagram from last lecture

Static	Dynamic
<ul style="list-style-type: none"> - set before execution - can be optimized up front - e.g. default in Tensorflow 1 	<ul style="list-style-type: none"> - set during forward pass - can be changed dynamically during execution - e.g. in PyTorch & Tensorflow 2



Deep Learning Package Design Choices

High-level programming language:


- Lua (Torch)
 - Python (Theano, Tensorflow, PyTorch, JAX)
 - ...
-
- We choose **PyTorch** because it is easy to understand for Python users

PyTorch

<https://pytorch.org>

- PyTorch is
 - A deep learning library (but not only that)
 - Written in Python, but based on Torch that was written in Lua
 - Developed by Facebook AI research group
 - Open Source since February 2017 😊
 - Since then, steadily growing, especially in NLP research!

Number of academic papers including “PyTorch” vs.
“Tensorflow”



PyTorch

- PyTorch provides
 - Predefined functions to compute gradients, optimisers, ...
 - Simple, object-oriented interface
 - ONNX (Open Neural Network Exchange) support to easily use trained models in other frameworks (e.g. on mobile devices, on the web, ...)
 - **GPU support!**

Using GPUs

- Why GPU support?

→ Neural network operations are mostly matrix multiplications

→ Matrix multiplications can be parallelized very effectively!

→ GPUs excell at tasks where the same operation is done on many data points


$$BE = \begin{bmatrix} 8 & 1 & 2 \\ -5 & 6 & 7 \end{bmatrix} \begin{bmatrix} -5 & 1 \\ 0 & 2 \\ -11 & 7 \end{bmatrix}$$

$$= \begin{bmatrix} (8)(-5) + (1)(0) + (2)(-11) & (8)(1) + (1)(2) + (2)(7) \\ (-5)(-5) + (6)(0) + (7)(-11) & (-5)(1) + (6)(2) + (7)(7) \end{bmatrix}$$

$$= \begin{bmatrix} -62 & 24 \\ -52 & 56 \end{bmatrix}$$

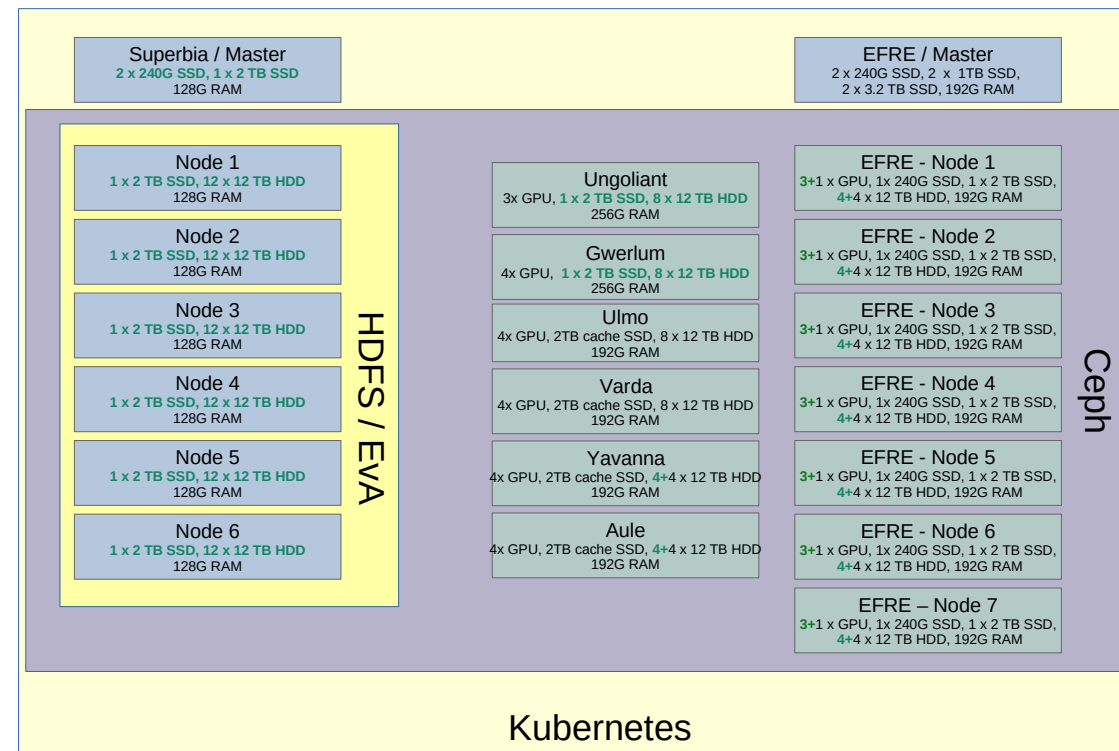
GPU vs CPU

- CPUs
 - Very few ($\sim 2 - 64$) complex cores
 - Great for complex tasks
 - Not so great for simple, parallel tasks
- GPUs
 - Have a lot ($\sim 2500 - 3800$) less complex cores
 - Can do huge matrix mults in parallel!



Deep Learning at our group

- Steadily growing server infrastructure
- **>60 GPUs** in >20 deep learning servers
- Additional university-wide GPU-cluster currently being built!
- Using Kubernetes to automatically manage resources



3.2 PyTorch

- How to use PyTorch?

From Numpy to PyTorch

- You may know Numpy
 - Math library for Python
 - Provides efficient implementation (C backend) for many common operations
- A simple program in Numpy:

```
In [23]: import numpy as np
```

```
In [24]: a = np.zeros((2,2)); b = np.ones((2,2))
```

```
In [25]: np.sum(b, axis=1)
```

```
Out[25]: array([ 2.,  2.])
```

```
In [26]: a.shape
```

```
Out[26]: (2, 2)
```

```
In [27]: np.reshape(a, (1,4))
```

```
Out[27]: array([[ 0.,  0.,  0.,  0.]])
```

From Numpy to PyTorch

- PyTorch can be seen as a Numpy replacement with GPU support
- Therefore very similar concepts and code:

```
import numpy as np
```

```
a = np.array([[1, 2], [3, 4]])  
b = np.ones((2, 2))
```

```
np.sum(b, axis=1)  
# array([2., 2.])
```

```
a.shape  
# (2, 2)
```

```
np.reshape(a, (1, 4))  
# array([[1, 2, 3, 4]])
```

possible to extract a
numpy array from a torch
tensor

```
import torch
```

```
a = torch.tensor([[1, 2], [3, 4]])  
b = torch.ones((2, 2))
```

```
torch.sum(b, dim=1)  
# tensor([2., 2.])
```


```
a.shape  
# torch.Size([2, 2])
```

```
torch.reshape(a, (1, 4))  
# tensor([[1, 2, 3, 4]])
```

```
a.numpy()  
# array([[1, 2],  
#        [3, 4]])
```





Neural Networks with PyTorch




$$L_{\theta}(l, x) = - \sum_i l[i] \cdot \log y_{\theta}(x)[i]$$

$$y = \text{softmax}(W_3 \tanh(W_2 \tanh(W_1 x + b_1) + b_2) + b_3)$$

$$x += -\text{learningrate} \cdot \frac{df}{dx}(x)$$



Neural Networks with PyTorch




- How to do this in PyTorch?

→ Three modules corresponding to these steps:

- nn Module
- Autograd module
- Optim module

PyTorch Modules



nn module

- Collection of neural network related building blocks, e.g.
 - fully connected layers
 - common activation functions
 - common loss functions
 - ...
- Handles weights and biases internally
→ hides complexity!

PyTorch Modules — Example



```

# Feedforward layer
torch.nn.Linear(...)
# Convolutional layer
torch.nn.Conv2d(...)
# ReLU activation
torch.nn.ReLU()
# Cross Entropy Loss
torch.nn.CrossEntropyLoss()
  
```

weights and bias →

- Tested code
- Object-oriented
- Automatically handles things like weight initialisation
- Easy to switch to other layers

```

[docs]@weak_module
class Linear(Module):
    r"""Applies a linear transformation to the incoming data:  $y = xA^T + b$ 

    Args:
        in_features: size of each input sample
        out_features: size of each output sample
        bias: If set to False, the layer will not learn an additive bias.
              Default: ``True``

    Shape:
        - Input:  $(N, *, \text{in\_features})$  where  $*$  means any number of
          additional dimensions
        - Output:  $(N, *, \text{out\_features})$  where all but the last dimension
          are the same shape as the input.

    Attributes:
        weight: the learnable weights of the module of shape
                 $(\text{out\_features}, \text{in\_features})$ . The values are
                initialized from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$ , where
                 $k = \frac{1}{\text{in\_features}}$ 
        bias: the learnable bias of the module of shape  $(\text{out\_features})$ .
              If attr: 'bias' is ``True``, the values are initialized from
               $\mathcal{U}(-\sqrt{k}, \sqrt{k})$  where
               $k = \frac{1}{\text{out\_features}}$ 

    Examples::

        >>> m = nn.Linear(20, 30)
        >>> input = torch.randn(128, 20)
        >>> output = m(input)
        >>> print(output.size())
        torch.Size([128, 30])

    """
    __constants__ = ['bias']


    def __init__(self, in_features, out_features, bias=True):
        super(Linear, self).__init__()
        self.in_features = in_features
        self.out_features = out_features
        self.weight = Parameter(torch.Tensor(out_features, in_features))
        if bias:
            self.bias = Parameter(torch.Tensor(out_features))
        else:
            self.register_parameter('bias', None)
        self.reset_parameters()

    def reset_parameters(self):
        init.kaiming_uniform_(self.weight, a=math.sqrt(5))
        if self.bias is not None:
            fan_in, _ = init._calculate_fan_in_and_fan_out(self.weight)
            bound = 1 / math.sqrt(fan_in)
            init.uniform_(self.bias, -bound, bound)

    @weak_script_method
    def forward(self, input):
        return F.linear(input, self.weight, self.bias)

    def extra_repr(self):
        return 'in_features={}, out_features={}, bias={}'.format(
            self.in_features, self.out_features, self.bias is not None
        )
  
```


PyTorch Modules



Autograd module

- Gradient computation necessary for gradient descent
 - Autograd abstracts backpropagation away
 - Builds computation graph
 - Performs differentiation by chain rule
- No differentiation by hand 😊

PyTorch Modules — Example



```
import torch
```

```
a = torch.tensor([[1,2],[3,4]], requires_grad=True)
```

```
# 1  2
```

```
# 3  4
```

```
y = torch.sum(a**2)
```

```
# => y = 1 + 4 + 9 + 16 = 30
```

We want to calculate
gradients w.r.t. this
tensor

```
# compute gradients of y w.r.t. to all Variables
```

```
y.backward()
```

```
a.grad # contains the gradient of y w.r.t. every entry of a
```


```
# 2  4
```

```
# 6  8
```

$$a = \begin{bmatrix} a_{00} & a_{01} \\ a_{10} & a_{11} \end{bmatrix}$$

$$y = \sum_{a_{ij}} a_{ij}^2 = a_{00}^2 + a_{10}^2 + a_{01}^2 + a_{11}^2$$

PyTorch Modules



Optim module

- Defines multiple optimisation algorithms
 - SGD
 - Adam optimiser
 - ...
- Consistent interface → easily interchangeable

These modules make it easy to build neural networks. Let's do this!

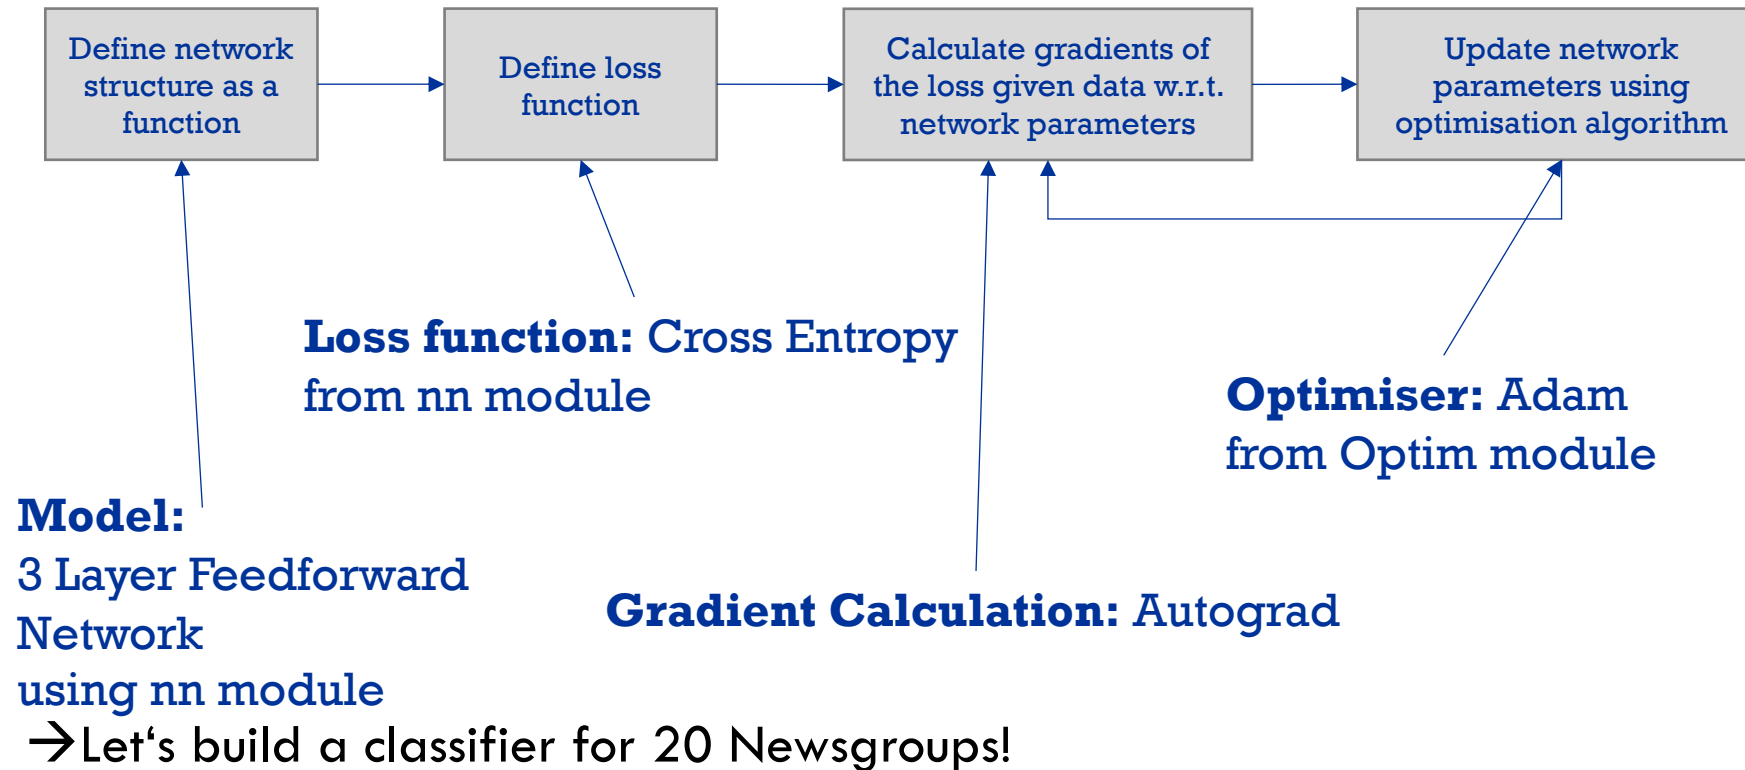


20 Newsgroups Text Classification

- 20 Newsgroups is a dataset of online discussion
 - 18,828 documents total
 - 20 forums, some closely related (pc.hardware vs mac.hardware), some highly unrelated (misc.forsale vs religion.christian)
- Classify which forum each document originates from
- Popular evaluation set in Natural Language Processing
- Not completely solved (Error rate $\sim 11.4\%$)



PyTorch – Building a Neural Network



But first: How do we get the data into the network?

Data Input and Preprocessing

PyTorch – Dataset

- Given:
 - A dataset of documents
 - For each document i
 - An attribute „data“ containing the raw unprocessed text
 - An attribute „topic“ containing the label
- We will build a Dataset class for 20 Newsgroups

```
{“topic”: “comp.os.ms-windows.misc”,  
 “data”: “I often use Notepad to view and  
 print "read.me" type files. I often  
 find\n> myself rushing to get to  
 Print Manager to stop the printer  
 and delete...”}
```



PyTorch – Dataset class

- PyTorch provides an abstract class representing a dataset
- Our Dataset class should inherit from it and overwrite the following methods:
 - `__len__` returns the length of our dataset
 - `__getitem__` returns a data point given an index
- `__getitem__` can load a specific data point on demand
→ no need to load the entire dataset

```
import numpy as np
from torch.utils.data import Dataset
```

```
class NewsgroupsDataset(Dataset):
```

```
    """20 Newsgroups Dataset"""
```

```
    def __init__(self, data: list, labels: dict):
```

initialise everything

```
    def __len__(self):
```

```
        """Returns the size of the dataset""",
```

get the number of examples

```
    def __getitem__(self, idx: int):
```

```
        """Returns a data point (text and label) given an index"""
```

get an example and label
for a given index

```
import numpy as np
from torch.utils.data import Dataset
```

```
class NewsgroupsDataset(Dataset):
```

```
    """20 Newsgroups Dataset"""
```

```
    def __init__(self, data: list, labels: dict):
        super().__init__()
        self.data = data # [{"data": "Hello there ...", "label": "misc.forsale"}, {"data": "...", "label": "..."}, ...]
        self.labels = labels # {"misc.forsale": 0, "sci.space": 1, ...}
```

initialise everything

```
    def __len__(self):
        """Returns the size of the dataset",,
```

get the number of examples

```
    def __getitem__(self, idx: int):
        """Returns a data point (text and label) given an index"""
```

get an example and label
for a given index

```
import numpy as np
from torch.utils.data import Dataset
```

```
class NewsgroupsDataset(Dataset):
```

```
    """20 Newsgroups Dataset"""
```

```
    def __init__(self, data: list, labels: dict):
        super().__init__()
        self.data = data # [{"data": "Hello there ...", "label": "misc.forsale"}, {"data": "...", "label": "..."}, ...]
        self.labels = labels # {"misc.forsale": 0, "sci.space": 1, ...}
```

initialise everything

```
    def __len__(self):
        """Returns the size of the dataset"""
        return len(self.data)
```

get the number of examples

```
    def __getitem__(self, idx: int):
        """Returns a data point (text and label) given an index"""
```

get an example and label
for a given index

```
import numpy as np
from torch.utils.data import Dataset
```

```
class NewsgroupsDataset(Dataset):
```

```
    """20 Newsgroups Dataset"""
```

```
    def __init__(self, data: list, labels: dict):
        super().__init__()
        self.data = data # [{"data": "Hello there ...", "label": "misc.forsale"}, {"data": "...", "label": "..."}, ...]
        self.labels = labels # {"misc.forsale": 0, "sci.space": 1, ...}
```

initialise everything

```
    def __len__(self):
        """Returns the size of the dataset"""
        return len(self.data)
```

get the number of examples

```
    def __getitem__(self, idx: int):
        """Returns a data point (text and label) given an index"""
        text = self.data[idx]["data"] # load the raw text from the document with the given id
        text = self.preprocess(text) # clean text (e.g. lowercasing, ...) and create embedding vector
        text = torch.from_numpy(text).float() # network inputs need to be float

        label = self.data[idx]["label"] # load the true label of the document with the given id
        label = self.labels[label] # lookup integer value of the string label
        label = torch.tensor(label).long() # label is not a continuous value but class indices

        return text, label
```

get an example and label
for a given index

PyTorch – Dataset

- Now we can retrieve an element using an index from 0 to $\text{len}(\text{dataset}) - 1$
 - One more thing:
Neural Networks use *batches* for training:
 - Concatenate multiple examples to a higher dimensional matrix
 - Train on multiple examples at once
- This is handled by a DataLoader in PyTorch
- Using a DataLoader also allows for:
 - Shuffling: Shuffle a list of all available indices and iterate over it
 - Parallel loading of items: Query the `__getitem__` method from multiple threads

```
# gensim provides a module for downloading datasets/models
import gensim.downloader as api
from torch.utils.data import DataLoader
```

```
data = list()
label_mask = set()
for document in api.load("20-newsgroups"):
    data.append(document)
    label_mask.add(document["topic"])
# assign an unique integer to every string label
label_mask = {label: index for index, label in enumerate(label_mask)}
```

```
dataset = NewsgroupsDataset(data, label_mask)
```

create a DataLoader by
handing in our dataset


```
data_loader = torch.utils.data.DataLoader(dataset, batch_size=512, shuffle=True,
num_workers=2)
```

```
for data in data_loader:
    inputs, labels = data
```

iterable that yields data batches
and their labels that can be used in
model training and testing

```
# feed inputs through network
# calculate loss based on output and labels
# ...
```

Building the Model




Two ways:

- The **hard/low-level** (but educational 🤔) way
- The **easy/high-level** way

PyTorch – Building the Model

- PyTorch provides an abstract class `Module`
(nearly everything that modifies tensors is a `Module`)
- Mainly two methods to override:
 - `forward`: feed incoming values through the network
 - `backward`: defines the gradient of the operation → Autograd handles this for us (when only using PyTorch operations), so usually no need to implement this
- We will build a simple feedforward neural network
 - With 2 hidden layers
 - ReLU activation
 - And a softmax output layer



Building Models: The **Hard/Low-Level** Way

Useful when:

- PyTorch does not (yet) provide layers you need
- you want to do other kinds of calculations

Low-Level Model Building

- Recall: A simple feedforward layer can be written as

$$y = \sigma(Wx + b)$$

- We can use PyTorch's functions that are inspired by Numpy to implement this!
- We need:
 - A matrix W
 - A bias term b
 - Operations
 - matrix multiplication
 - addition
 - element-wise operations for the activation function

Low-Level Model Building

- **One addition:** Up until now, we always used *column vectors*
- PyTorch, however, uses *row vectors*
- Therefore, the linear classifier

$$y = \sigma(Wx + b)$$

needs to be rewritten as

$$y = \sigma(xW + b)$$

to be calculable.

```
class NewsgroupsModelLowLevel(nn.Module):
    """Simple Feedforward Neural Network for 20 Newsgroups"""
```

```
def __init__(self, input_size=300):
    super().__init__()
    self.input_size = input_size
    self.hidden_1_size = 2048
    self.hidden_2_size = 256
    self.num_classes = 20
```

weights and biases are randomly initialised

```
self.W1 = nn.Parameter(torch.randn(self.input_size, self.hidden_1_size, requires_grad=True))
self.b1 = nn.Parameter(torch.randn(1, self.hidden_1_size, requires_grad=True))
self.relu1 = nn.ReLU()

self.W2 = nn.Parameter(torch.randn(self.hidden_1_size, self.hidden_2_size, requires_grad=True))
self.b2 = nn.Parameter(torch.randn(1, self.hidden_2_size, requires_grad=True))
self.relu2 = nn.ReLU()

self.W3 = nn.Parameter(torch.randn(self.hidden_2_size, self.num_classes, requires_grad=True))
self.b3 = nn.Parameter(torch.randn(1, self.num_classes, requires_grad=True))
```

a {
b {
c {


```
def forward(self, x):
    # first hidden layer
    a = x @ self.W1 + self.b1
    a = self.relu1(a)
    # second hidden layer
    b = a @ self.W2 + self.b2
    b = self.relu2(b)
    # output layer
    c = b @ self.W3 + self.b3
    return c # => logits
```

wrapped by Parameter to be later
picked up by the optimisation

data input

no softmax

@ is the shorthand for matrix multiplication



Building Models: The **Easy/High-Level** Way

Most layers are already implemented in PyTorch's nn Module!

```
import torch.nn as nn
```

```
class NewsgroupsModel(nn.Module):
```

```
    """Simple Feedforward Neural Network for 20 Newsgroups"""
```

```
    def __init__(self, input_size=300):
```

```
        super().__init__()
```

```
        self.image_size = input_size
```

```
        self.hidden_1_size = 2048
```

takes care of weights, biases, initialisation, ... a

```
        self.hidden_2_size = 256
```

```
        self.num_classes = 20
```

```
        self.fc1 = nn.Linear(self.image_size, self.hidden_1_size)
```

```
        self.relu1 = nn.ReLU()
```

```
        self.fc2 = nn.Linear(self.hidden_1_size, self.hidden_2_size)
```

```
        self.relu2 = nn.ReLU()
```

```
        self.fc3 = nn.Linear(self.hidden_2_size, self.num_classes)
```

```
    def forward(self, x):
```

```
        a = self.relu1(self.fc1(x))
```


```
        b = self.relu2(self.fc2(a))
```

```
        c = self.fc3(b) # => logits
```

```
    return c
```

data input

still no softmax



Short slide-in: Initialising Neural Networks

- The weights in a network need some initial values
 - PyTorch layers handle this for us
 - But what is the best option for initialisation?
 - **Our low-level approach was:** Take a normal distribution around zero and pick some values...
 - ... was that a smart idea?
- Maybe. It depends

Initialising Neural Networks

- Many possible initialisers
 - Problem: You never know which works best in your context
- Try different initialisers
- Some common choices on the following slides

Initialising Neural Networks – Random

- The easiest initialiser: Just sample random values
- Two variants:
 - Pick numbers from a **uniform** distribution, usually close to zero
`w = np.random.uniform(-0.01, 0.01)`
 - Pick numbers from a **normal** distribution, usually around zero
`w = 0.01 * np.random.randn()`

Initialising Neural Networks – Calibrating the Variance

- Problem with purely random initialisation:
More inputs
→ Higher (variance of the) output of the neuron
→ Possible problem with high gradients!
- Can be fixed by „normalising“ the initialisation
- For each weight w of a neuron N :
 $w = \text{np.random.randn}() / \text{sqrt}(n)$
- n is the number of inputs to the neuron N

Initialising Neural Networks – Glorot/Xavier

- Introduced by Xavier Glorot and Yoshua Bengio
- Names **Glorot-** and **Xavier-Initialiser** used interchangeably
- Complex analysis of gradient flow in networks
→ Recommend to normalise the variance to

$$Var(w) = \frac{2}{n_{in} + n_{out}}$$

Number of inputs/outputs of the neuron

```
w = np.random.randn() / sqrt(2 / (n_in + n_out))
```

- Works best for layers with sigmoid or tanh activation functions

Initialising Neural Networks – He/Kaiming

- Introduced by Kaiming He et al.
- Names **He-** and **Kaiming-Initialiser** used interchangeably

$$Var(w) = \frac{gain}{n_{in}}$$

- *gain* depends on activation function, e.g. for ReLU: *gain* = 2

```
w = np.random.randn() / sqrt(2/n_in)
```

- Was initially introduced for the ReLU activation function, thus works best for it

HE, Kaiming, et al. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. In: *Proceedings of the IEEE international conference on computer vision*. 2015. S. 1026-1034.

Initialising Neural Networks - PyTorch

- PyTorch contains many common initialisers
 - Random uniform/normal
 - Glorot/Xavier
 - He/Kaiming
 - ...
 - He/Kaiming initialisation is the default for most layers
- Use He/Kaiming first. Try others, too!

Initialising Neural Networks - PyTorch

How to initialise weights in PyTorch?

1. Initialise only one layer after creating it:


```
torch.nn.init.xavier_uniform(layer.weight)
```

2. Initialise the whole model after creating it

```
def init_weights(m):  
    if type(m) == nn.Linear:  
        torch.nn.init.xavier_uniform(m.weight)
```

```
model.apply(init_weights)
```

Loss Function



PyTorch – Loss Function

- Given the output of the network, calculate the error of the prediction w.r.t. the correct label → Loss function / „criterion“

```
import torch.nn as nn
```

 ← nn module provides common loss functions

```
# Loss function
```

```
criterion = nn.CrossEntropyLoss()
```

applies softmax for us as gradient calculation is faster this way


```
# calculate the loss
```

```
loss = criterion(logits, labels)
```

Unnormalised output of the network

Correct label

Gradient Calculation



PyTorch – Gradient Calculation

- Given the value of the loss function and the current weights, calculate the gradient for all parameters

```
import torch.nn as nn
```

```
# Loss function
```

```
criterion = nn.CrossEntropyLoss()
```

```
# calculate the loss
```

```
loss = criterion(logits, labels)
```


```
[...]
```

```
# calculate the gradients w.r.t. the network parameters
```

```
loss.backward()
```

That was easy 😊

Optimiser



PyTorch – Optimiser

- Alter the weights of the network to minimise the error
→ Optimiser

optim provides common optimisers

```
import torch.optim as optim
```

learning rate


```
# Optimiser
```

```
optimiser = optim.Adam(model.parameters(), lr=0.001)
```

all weights and biases of the NewsgroupsModel instance

That was also easy 😊

Putting things together: The Training Loop



PyTorch — Training Loop

```
for epoch in range(10): # loop over the dataset multiple times
    for data in data_loader:
        # get the data points
        inputs, labels = data

        # zero the parameter gradients
        # (else, they are accumulated)
        optimiser.zero_grad()

        # forward the data through the network
        logits = model(inputs)

        # calculate the loss
        loss = criterion(logits, labels)

        # calculate the gradients w.r.t. the network parameters
        loss.backward()

        # let the optimiser take an optimization step
        optimiser.step()
```

```
# Data
dataset = NewsgroupsDataset(data)
data_loader =
torch.utils.data.DataLoader(datas
et, batch_size=512, shuffle=True,
num_workers=2)

# Model
model = NewsgroupsModel()

# Loss function
criterion = nn.CrossEntropyLoss()

# Optimiser
optimiser =
optim.Adam(model.parameters(),
lr=0.001)
```

PyTorch – A Model for 20 Newsgroups

We've set up everything, it's time to let it run

→ Demo Time!

Combating Overfitting – Regularisation

- Our network is adequate on the training set (78.96 % Accuracy)...
- ... but worse on the test set (70.17 % Accuracy)!
- **Overfitting** is very common for neural networks
- Deep networks → Lots of parameters → Memorising the training set
- **Regularisation** is designed to prevent this
- Two main tactics:
 - Penalising high weights (L1 / L2-norm)
 - Dropout

Combating Overfitting – L2 Regularisation

- **Idea:**
Force the network to use all inputs rather than focusing on some
- Modify the loss to penalise „peaky“ weights
- Add a regularisation term:

$$L_{reg} = L + \frac{1}{2}\lambda \sum_w w^2$$

Loss without regularisation

Regularisation „strength“

All weights in the network

- **L_2 -norm** is smaller if there are no „outliers“ in the weights

Combating Overfitting – L2 Regularisation

- L2-Regularisation in PyTorch:

$$L_{reg} = L + \frac{1}{2} \lambda \sum_w w^2$$

- You could add the loss term by hand:

```
for param in model.parameters():  
    loss += 0.5 * lamb * torch.sum(param**2)
```

weight or bias tensor from the defined
layers

- But most optimisers already have a `weight_decay` parameter that is closely related and mostly equivalent to L2 regularisation:

```
optimiser = optim.Adam(model.parameters(), lr=0.001, weight_decay=lamb)
```

More information about the similarities and differences between L2 regularisation and weight decay can be found at <https://bbabenko.github.io/weight-decay/>

L2 Regularisation vs. Weight Decay

- L2 Regularisation adds a term to **the loss function**:

$$L_{reg} = L + \underbrace{\frac{1}{2}\lambda \sum_w w^2}_{\text{that's new}}$$

- Deriving the regularised loss function:

$$\frac{\partial}{\partial w} L_{reg} = \frac{\partial}{\partial w} L + \lambda w$$

- Weight decay adds a term to **the derivative of the loss function**:

$$\frac{\partial}{\partial w} L + \underbrace{\lambda w}_{\text{that's new}}$$

- Then, in the optimisation step (SGD):

$$w = w - lr \cdot \left(\frac{\partial}{\partial w} L + \lambda w \right)$$

looks like they are equivalent!

L2 Regularisation vs. Weight Decay


Question: But are they? When do they obtain different results?

Answer:


- Adding a term to the optimisation step (**Weight Decay**)
vs.
modifying the loss gradient (**L2**)
- Equivalent for optimisers that do not reuse previous gradients (e.g. SGD)
- Not equivalent for optimisers that reuse previous loss gradients (e.g. Adam or SGD with Momentum)
 - L2 affects the previous gradient
 - The Weight Decay term is not reused

Combating Overfitting – Dropout

- Dropout: More recent method (2014)
- **Idea:** During **training**, „remove“ a portion $0 \leq p < 1$ of neurons from the net
→ Force the net to rely on many features instead of few!



(a) Standard Neural Net



(b) After applying dropout.

Combating Overfitting – Dropout

- Implementation by a **mask**:

Random vector with zeros (probability **p**) and ones

Fully Connected layer with ReLU

```
H1 = np.maximum(0, np.dot(W1, X) + b1)
```

```
U1 = np.random.rand(*H1.shape) > p
```

→ Mask with zeros and ones

```
H1 *= U1
```

→ Keep the neuron if the position in the mask has a one, set it to zero otherwise

→ Equivalent to removing the neuron


Combating Overfitting – Dropout

- Dropout only applied in **training** phase!
 - When **testing**, dropout is **disabled** ($p = 0$)
- Roughly equivalent to training an **ensemble** of smaller networks
- Empirically shown to improve generalisation on many tasks

Combating Overfitting – Dropout



Problem: Dropout changes the expected output of a neuron!



- Larger expected output when testing!
- Scaling needed!

Combating Overfitting – Dropout

- Scaling needed because of different expected output during training/testing


- Two possible solutions:

- Scale **down** while **testing**

```
H1 = np.maximum(0, np.dot(W1, X) + b1) * p
```

- Scale **up** while **training**

```
H1 = np.maximum(0, np.dot(W1, X) + b1)  
U1 = (np.random.rand(*H1.shape) > p) / p  
H1 *= U1
```



Combating Overfitting – Dropout

- Scaling needed because of different expected output during training/testing

- Two possible solutions:

- Scale **down** while **testing**

```
H1 = np.maximum(0, np.dot(W1, X) + b1)
```

This slows down the prediction step! We **don't** want that!

- Scale **up** while **training**

```
H1 = np.maximum(0, np.dot(W1, X) + b1)  
U1 = (np.random.rand(*H1.shape) > p) / p  
H1 *= U1
```

Combating Overfitting – Dropout

- PyTorch provides a Dropout module

```
# Model __init__  
self.dropout = nn.Dropout(p=0.4)
```

```
# Model forward  
x = self.dropout(x)
```

to drop values with a **probability** of 40% during training.

Important:

```
model.train() # default  
model.eval()
```

sets a variable in the model to indicate training model,
so dropout should be applied

sets the model training variable to False, so dropout is disabled;
call this before testing!



Attention!
Libraries use **keep probability** and **drop probability** inconsistently!

PyTorch – A Model for 20 Newsgroups

Dropout Demo

PyTorch – A Model for 20 Newsgroups

→ Better results on test set (71.35 % vs 70.17 % without Dropout)!

This may not sound like much, but:

- Additional **>1%** sometimes means a huge gain!
- Often bigger improvements on other datasets...
- ... all the more so if we consider the small effort involved.

- Curious effect of Dropout:
 - Results on the training set sometimes worse than on the test set
 - Caused by aforementioned „ensemble“ of smaller networks:
All neurons available for testing, only some for training
- Nice to have:
 - Dropout usually makes training faster due to more multiplications with zeros

PyTorch – Bringing everything to the GPU

- Currently, everything was computed on the CPU
- In PyTorch, we have to move all values to the GPU if we want to work with it
- Only a few changes are necessary to run on the GPU

```
# Automatically choose GPU if available  
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
# Move the data tensors (inputs and outputs) to the correct device  
input_tensor.to(device)  
output_tensor.to(device)
```

```
# Move the model parameters to the correct device  
model.to(device)
```

Making PyTorch development easier/faster

ECOSYSTEM TOOLS

pytorch.org/ecosystem/

PyTorch Lightning

PyTorch Lightning is a Keras-like ML library for PyTorch. It leaves core training and validation logic to you and automates the rest.

fastai

fastai is a library that simplifies training fast and accurate neural nets using modern best practices.

Ignite

Ignite is a high-level library for training neural networks in PyTorch. It helps with writing compact, but full-featured training loops.

Pyro

Pyro is a universal probabilistic programming language (PPL) written in Python and supported by PyTorch on the backend.

... and many more

3.3 Vectorisation of Neural Networks

Why Vectorisation?

- **Recall:**

- Matrix multiplications can be computed very efficiently
- Element-wise operations can be parallelised (good for GPUs)


- **Idea:** Make your network faster to compute by using as many matrix multiplications and element-wise operations as possible

→ Vectorise the forward pass and backpropagation algorithm



Vectorising the Forward Pass

Vectorising the Forward Pass




You
already
know this

- Two steps:

- $z = \sum_{i=1}^3 w_i x_i$

- $y = \sigma(z)$

Vectorising the Forward Pass



Instead of calculating


$$z = \sum_{i=1}^3 w_i x_i$$

using a for-loop, we can compute

$$z = w^T x$$

using vector/matrix multiplication.

Vectorising the Forward Pass




- $\sigma(z)$ is just a function working on one scalar

→ Forward pass is easily parallelisable for one neuron 😊


→ But wait! Usually, we have more than one neuron per layer!

Vectorising the Forward Pass



- Fortunately, all neurons are the same, except the weights
- Combine all weight vectors w to matrix W :
 $W_{i,j}$ is the weight for the connection from neuron x_j to z_i

Vectorising the Forward Pass



Again two steps:

this is now a vector

$$z = Wx$$

this is now a matrix

$$h = \sigma(z)$$

element-wise
operation
(parallelisable!)

this is now a
vector

very similar to before,
but fully vectorised! 😊


Vectorising Backpropagation

Backpropagation

- Recall: Backpropagation =
 - calculating **partial derivatives**...
 - ... of the network's **loss function**...
 - ... w.r.t. the **weights and biases**...
 - ... by applying the **chain rule**:


$$\text{If } f(x) = p(q(x)), \text{ then } \frac{df}{dx} = \frac{df}{dq} \frac{dq}{dx}$$

Multivariable Chain Rule



- In neural networks, it is typical for a value to „flow“ through multiple paths to a neuron (here: $x \rightarrow y$)
- When calculating gradients w.r.t. x , we can go two paths:
 - via z_1 : $\frac{dy}{dz_1} \frac{dz_1}{dx}$
 - via z_2 : $\frac{dy}{dz_2} \frac{dz_2}{dx}$
- How do we calculate the final gradient $\frac{dy}{dx}$?

Multivariable Chain Rule



- The multivariable chain rule states:

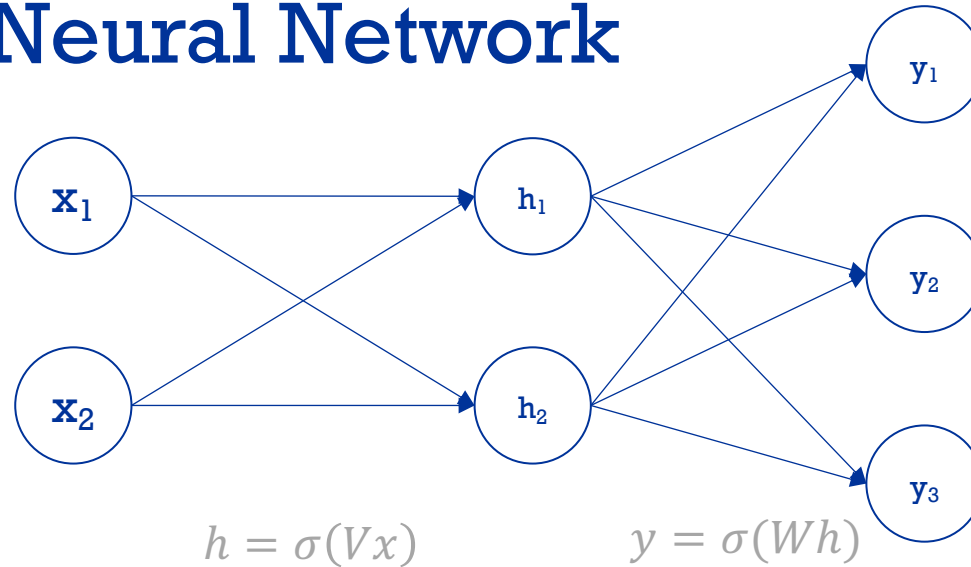
$$\text{If } f(g(x), q(x)), \text{ then } \frac{df}{dx} = \frac{df}{dg} \frac{dg}{dx} + \frac{df}{dq} \frac{dq}{dx}$$

- In this case:

$$\frac{dy}{dx} = \underbrace{\frac{dy}{dz_1} \frac{dz_1}{dx}}_{\text{path via } z_1} + \underbrace{\frac{dy}{dz_2} \frac{dz_2}{dx}}_{\text{path via } z_2}$$

Okay, let's vectorise the backpropagation in
an example neural network!

Example Neural Network





$$L(y, t) = \frac{1}{2} \sum_{i=1}^3 (t_i - y_i)^2$$

output vector target vector

- Simple Feedforward Network with one hidden layer
- No bias, but possible to add using the bias trick (see exercises)
- Sigmoid activation function
- Loss function: squared error loss


First, let's visualise this network as a circuit diagram with vectors

Example Neural Network



Backpropagation: Use the chain rule and go back every step in the diagram

Backpropagation Step 1



what is $\frac{\partial l}{\partial y}$?

$$L(y, t) = \frac{1}{2} \sum_{i=1}^3 (t_i - y_i)^2$$


$$\frac{\partial l}{\partial y_1} = \frac{\partial}{\partial y_1} \frac{1}{2} \sum_{i=1}^3 (t_i - y_i)^2$$

$$\frac{\partial l}{\partial y_2} = \frac{\partial}{\partial y_2} \frac{1}{2} \sum_{i=1}^3 (t_i - y_i)^2$$

$$\frac{\partial l}{\partial y_3} = \frac{\partial}{\partial y_3} \frac{1}{2} \sum_{i=1}^3 (t_i - y_i)^2$$

$$\frac{\partial l}{\partial y} = \begin{bmatrix} \frac{\partial l}{\partial y_1} \\ \frac{\partial l}{\partial y_2} \\ \frac{\partial l}{\partial y_3} \end{bmatrix}$$

Backpropagation Step 1



$$\frac{\partial l}{\partial y_1} = \frac{\partial}{\partial y_1} \frac{1}{2} \sum_{i=1}^3 (t_i - y_i)^2$$

$$= \frac{\partial}{\partial y_1} \left(\frac{1}{2} (t_1 - y_1)^2 + \frac{1}{2} (t_2 - y_2)^2 + \frac{1}{2} (t_3 - y_3)^2 \right)$$


sum rule

$$= \frac{\partial}{\partial y_1} \frac{1}{2} (t_1 - y_1)^2 + \cancel{\frac{\partial}{\partial y_1} \frac{1}{2} (t_2 - y_2)^2} + \cancel{\frac{\partial}{\partial y_1} \frac{1}{2} (t_3 - y_3)^2}$$

$$= \frac{\partial}{\partial y_1} \boxed{\frac{1}{2} (t_1 - y_1)^2}$$

chain rule: $p(q(y_1))$ with
 $\boxed{q(y_1, t_1) = t_1 - y_1}$
 $\boxed{\text{and } p(q) = \frac{1}{2} (q)^2}$

Backpropagation Step 1




chain rule: $p(q(y_1, t_1))$ with
 $q(y_1, t_1) = t_1 - y_1$
 and $p(q) = \frac{1}{2}(q)^2$

Previous slide

$$\frac{\partial l}{\partial y_1} = \frac{\partial}{\partial y_1} \frac{1}{2} (t_1 - y_1)^2 = \underbrace{2 \frac{1}{2} (t_1 - y_1)}_{\frac{dp}{dq}} \cdot \underbrace{-1}_{\frac{dq}{dy_1}} = y_1 - t_1$$

Backpropagation Step 1



$$\frac{\partial l}{\partial y_1} = y_1 - t_1$$


$$\frac{\partial l}{\partial y_2} = y_2 - t_2$$

$$\frac{\partial l}{\partial y_3} = y_3 - t_3$$

$$\frac{\partial l}{\partial y} = \begin{bmatrix} \frac{\partial l}{\partial y_1} \\ \frac{\partial l}{\partial y_2} \\ \frac{\partial l}{\partial y_3} \end{bmatrix} = \boxed{y - t}$$

fast vector operation 😊

Backpropagation Step 2



σ is an element-wise operation, so we only need to look at the i^{th} element

chain rule!


$$\frac{\partial l}{\partial b} = \begin{bmatrix} \frac{\partial l}{\partial b_1} \\ \frac{\partial l}{\partial b_2} \\ \frac{\partial l}{\partial b_3} \end{bmatrix}$$

$$\frac{\partial l}{\partial b_1} = \frac{\partial l}{\partial y_1} \cdot \frac{\partial y_1}{\partial b_1} = (y_1 - t_1) \cdot \sigma(b_1) \cdot (1 - \sigma(b_1))$$

$$\frac{\partial l}{\partial b_2} = \frac{\partial l}{\partial y_2} \cdot \frac{\partial y_2}{\partial b_2} = (y_2 - t_2) \cdot \sigma(b_2) \cdot (1 - \sigma(b_2))$$

$$\frac{\partial l}{\partial b_3} = \frac{\partial l}{\partial y_3} \cdot \frac{\partial y_3}{\partial b_3} = (y_3 - t_3) \cdot \sigma(b_3) \cdot (1 - \sigma(b_3))$$

Backpropagation Step 3




what are $\frac{\partial l}{\partial h}$ and $\frac{\partial l}{\partial W}$?

$$\frac{\partial l}{\partial h} = \begin{bmatrix} \frac{\partial l}{\partial h_1} \\ \frac{\partial l}{\partial h_2} \end{bmatrix}$$


We need this for calculating the derivatives of the previous layer

let's start with this

$$\frac{\partial l}{\partial W} = \begin{bmatrix} \frac{\partial l}{\partial W_{1,1}} & \frac{\partial l}{\partial W_{1,2}} \\ \frac{\partial l}{\partial W_{2,1}} & \frac{\partial l}{\partial W_{2,2}} \\ \frac{\partial l}{\partial W_{3,1}} & \frac{\partial l}{\partial W_{3,2}} \end{bmatrix}$$



Backpropagation Step 3 — h



$$\frac{\partial l}{\partial h} = \begin{bmatrix} \frac{\partial l}{\partial h_1} \\ \frac{\partial l}{\partial h_2} \end{bmatrix}$$

what is $\frac{\partial l}{\partial h}$?


$$b = Wh = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} = \begin{bmatrix} W_{1,1}h_1 + W_{1,2}h_2 \\ W_{2,1}h_1 + W_{2,2}h_2 \\ W_{3,1}h_1 + W_{3,2}h_2 \end{bmatrix}$$

multivariable chain rule!

$$\frac{\partial l}{\partial h_1} = \frac{\partial l}{\partial b_1} \cdot \frac{\partial b_1}{\partial h_1} + \frac{\partial l}{\partial b_2} \cdot \frac{\partial b_2}{\partial h_1} + \frac{\partial l}{\partial b_3} \cdot \frac{\partial b_3}{\partial h_1} = \frac{\partial l}{\partial b_1} \cdot W_{1,1} + \frac{\partial l}{\partial b_2} \cdot W_{2,1} + \frac{\partial l}{\partial b_3} \cdot W_{3,1}$$

$$\frac{\partial l}{\partial h_2} = \frac{\partial l}{\partial b_1} \cdot \frac{\partial b_1}{\partial h_2} + \frac{\partial l}{\partial b_2} \cdot \frac{\partial b_2}{\partial h_2} + \frac{\partial l}{\partial b_3} \cdot \frac{\partial b_3}{\partial h_2} = \frac{\partial l}{\partial b_1} \cdot W_{1,2} + \frac{\partial l}{\partial b_2} \cdot W_{2,2} + \frac{\partial l}{\partial b_3} \cdot W_{3,2}$$

Backpropagation Step 3 — h




From previous slide:

$$\frac{\partial l}{\partial h} = \begin{bmatrix} \frac{\partial l}{\partial h_1} \\ \frac{\partial l}{\partial h_2} \end{bmatrix} = \begin{bmatrix} \frac{\partial l}{\partial b_1} \cdot W_{1,1} + \frac{\partial l}{\partial b_2} \cdot W_{2,1} + \frac{\partial l}{\partial b_3} \cdot W_{3,1} \\ \frac{\partial l}{\partial b_1} \cdot W_{1,2} + \frac{\partial l}{\partial b_2} \cdot W_{2,2} + \frac{\partial l}{\partial b_3} \cdot W_{3,2} \end{bmatrix} = \boxed{W^T \frac{\partial l}{\partial b}} \quad \text{vectorised version 😊}$$

Backpropagation Step 3

You
already
know this




what are $\frac{\partial l}{\partial h}$ and $\frac{\partial l}{\partial W}$?

$$\frac{\partial l}{\partial h} = \begin{bmatrix} \frac{\partial l}{\partial h_1} \\ \frac{\partial l}{\partial h_2} \end{bmatrix}$$


We need this for calculating the derivatives of the previous layer

$$\frac{\partial l}{\partial W} = \begin{bmatrix} \frac{\partial l}{\partial W_{1,1}} & \frac{\partial l}{\partial W_{1,2}} \\ \frac{\partial l}{\partial W_{2,1}} & \frac{\partial l}{\partial W_{2,2}} \\ \frac{\partial l}{\partial W_{3,1}} & \frac{\partial l}{\partial W_{3,2}} \end{bmatrix}$$

now this one



Backpropagation Step 3 — W



$$\frac{\partial l}{\partial W} = \begin{bmatrix} \frac{\partial l}{\partial W_{1,1}} & \frac{\partial l}{\partial W_{1,2}} \\ \frac{\partial l}{\partial W_{2,1}} & \frac{\partial l}{\partial W_{2,2}} \\ \frac{\partial l}{\partial W_{3,1}} & \frac{\partial l}{\partial W_{3,2}} \end{bmatrix}$$

$$b = Wh = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} = \begin{bmatrix} W_{1,1}h_1 + W_{1,2}h_2 \\ W_{2,1}h_1 + W_{2,2}h_2 \\ W_{3,1}h_1 + W_{3,2}h_2 \end{bmatrix}$$


multivariable chain rule!

$$\frac{\partial l}{\partial W_{1,1}} = \frac{\partial l}{\partial b_1} \cdot \frac{\partial b_1}{\partial W_{1,1}} + \frac{\partial l}{\partial b_2} \cdot \frac{\partial b_2}{\partial W_{1,1}} + \frac{\partial l}{\partial b_3} \cdot \frac{\partial b_3}{\partial W_{1,1}} = \frac{\partial l}{\partial b_1} \cdot h_1 + \frac{\partial l}{\partial b_2} \cdot 0 + \frac{\partial l}{\partial b_3} \cdot 0$$

$$\vdots$$

$$\frac{\partial l}{\partial W_{3,2}} = \frac{\partial l}{\partial b_1} \cdot \frac{\partial b_1}{\partial W_{3,2}} + \frac{\partial l}{\partial b_2} \cdot \frac{\partial b_2}{\partial W_{3,2}} + \frac{\partial l}{\partial b_3} \cdot \frac{\partial b_3}{\partial W_{3,2}} = \frac{\partial l}{\partial b_1} \cdot 0 + \frac{\partial l}{\partial b_2} \cdot 0 + \frac{\partial l}{\partial b_3} \cdot h_2$$


Backpropagation Step 3 — W



$$\frac{\partial l}{\partial b} = \begin{bmatrix} \frac{\partial l}{\partial b_1} \\ \frac{\partial l}{\partial b_2} \\ \frac{\partial l}{\partial b_3} \end{bmatrix}$$

$$\frac{\partial l}{\partial W} = \begin{bmatrix} \frac{\partial l}{\partial W_{1,1}} & \frac{\partial l}{\partial W_{1,2}} \\ \frac{\partial l}{\partial W_{2,1}} & \frac{\partial l}{\partial W_{2,2}} \\ \frac{\partial l}{\partial W_{3,1}} & \frac{\partial l}{\partial W_{3,2}} \end{bmatrix} = \begin{bmatrix} \frac{\partial l}{\partial b_1} \cdot h_1 & \frac{\partial l}{\partial b_1} \cdot h_2 \\ \frac{\partial l}{\partial b_2} \cdot h_1 & \frac{\partial l}{\partial b_2} \cdot h_2 \\ \frac{\partial l}{\partial b_3} \cdot h_1 & \frac{\partial l}{\partial b_3} \cdot h_2 \end{bmatrix} = \boxed{\frac{\partial l}{\partial b} h^T} \text{ vectorised version 😊}$$

Backpropagation Step 4



what is $\frac{\partial l}{\partial a}$?

σ is an element-wise operation, so we only need to look at the i^{th} element


chain rule!

$$\frac{\partial l}{\partial a} = \begin{bmatrix} \frac{\partial l}{\partial a_1} \\ \frac{\partial l}{\partial a_2} \end{bmatrix}$$

$$\begin{aligned} \frac{\partial l}{\partial a_1} &= \frac{\partial l}{\partial h_1} \cdot \frac{\partial h_1}{\partial a_1} = \frac{\partial l}{\partial h_1} \cdot \sigma(a_1) \cdot (1 - \sigma(a_1)) \\ \frac{\partial l}{\partial a_2} &= \frac{\partial l}{\partial h_2} \cdot \frac{\partial h_2}{\partial a_2} = \frac{\partial l}{\partial h_2} \cdot \sigma(a_2) \cdot (1 - \sigma(a_2)) \end{aligned}$$

same structure as in Step 2, when we calculated $\frac{\partial l}{\partial W}$

Backpropagation Step 5




what are $\frac{\partial l}{\partial x}$ and $\frac{\partial l}{\partial V}$?

$$\frac{\partial l}{\partial x} = \begin{bmatrix} \frac{\partial l}{\partial x_1} \\ \frac{\partial l}{\partial x_2} \end{bmatrix}$$


We don't need these for weight optimisation (but we could)
→ Don't calculate them

$$\frac{\partial l}{\partial V} = \begin{bmatrix} \frac{\partial l}{\partial V_{1,1}} & \frac{\partial l}{\partial V_{1,2}} \\ \frac{\partial l}{\partial V_{2,1}} & \frac{\partial l}{\partial V_{2,2}} \end{bmatrix}$$



same structure as in Step 3, when we calculated $\frac{\partial l}{\partial W}$

Backpropagation Step 5 — V




Step 3 result:

$$\frac{\partial l}{\partial W} = \frac{\partial l}{\partial b} h^T \quad \text{with } b = Wh$$

Same calculation works here:

$$a = Vx$$

$$\frac{\partial l}{\partial V} = \begin{bmatrix} \frac{\partial l}{\partial V_{1,1}} & \frac{\partial l}{\partial V_{1,2}} \\ \frac{\partial l}{\partial V_{2,1}} & \frac{\partial l}{\partial V_{2,2}} \end{bmatrix} = \boxed{\frac{\partial l}{\partial a} x^T} \quad \text{vectorised calculation 😊}$$



Summary Vectorising Neural Networks

- How can we speed up neural network calculations?
 1. Backpropagation algorithm: reuse calculated derivatives
 2. Vectorisation: use highly optimised vector and matrix multiplications
 3. Faster hardware: GPUs, parallelize computations, ...

} made computations feasible in 1970s

} made computing big models possible (since 2009)
- Trends since then:
 - Faster hardware
 - More efficient models than feedforward networks (e.g. CNN, that reuses weights)
 - Optimisation algorithms that converge faster (less training iterations → faster training)

Next Week:

Convolutional Neural Networks!